

Blockchains & Distributed Ledgers

Lecture 02

Aggelos Kiayias

Slide credits: AK, Dionysis Zindros, Christos Nasikas

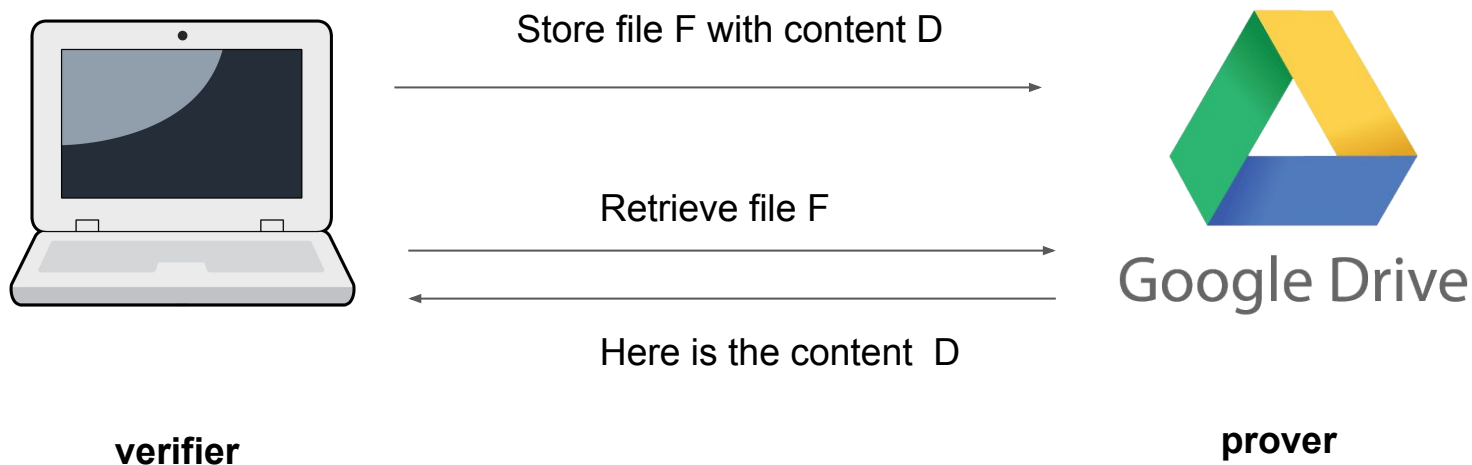
Overview

- Motivation: Server file storage
- Merkle trees to store lists
- Proofs-of-inclusion
- Merkle trees to store sets
- Proofs-of-non-inclusion
- Merkle–Patricia tries to store key:value pairs
- Blocks and blockchains

Authenticated Data Structures

- Like regular data structures, but cryptographically authenticated
- Allows a **verifier** to store, retrieve and operate on data with an untrusted **prover**

The authenticated file storage problem



The file storage problem

- Client wants to store a file on a server
- File has a name F and content D
- Clients wants to retrieve file F later

File storage: Basic protocol

- Client sends file F with content D to server
- Server stores (F, D)
- Client deletes D
- Client requests F from server
- Server returns D
- Client has recovered D

File storage: Protocol against adversaries

- What if **server is adversarial** and returns $D' \neq D$?

File storage: Protocol against adversaries

Trivial solution:

- Client does not delete D
- When server returns D' , client compares D and D'

...what if client doesn't have enough memory to store D for a long time?

File storage: Hash-based protocol


- Client sends file F with data D to server
- Server stores (F, D)
- Client stores $H(D)$, deletes D
- Client requests F from server
- Server returns D'
- Client compares $H(D') = H(D)$

File storage: File chunks

- What if client wants to retrieve the 200,019th byte of the file?
- Must download the whole file...
- Merkle trees to the rescue!


Merkle Tree

- An **authenticated** binary tree
- Split file into **chunks** of, say, 1KB




Merkle Tree

- **Hash** each chunk using a cryptographic hash function (SHA256)
- Convention: Arrows show direction of hash function application




Merkle Tree

- **Combine** them by two to create a binary tree
- Each node stores the **hash** of the **concat** of its children




Merkle Tree



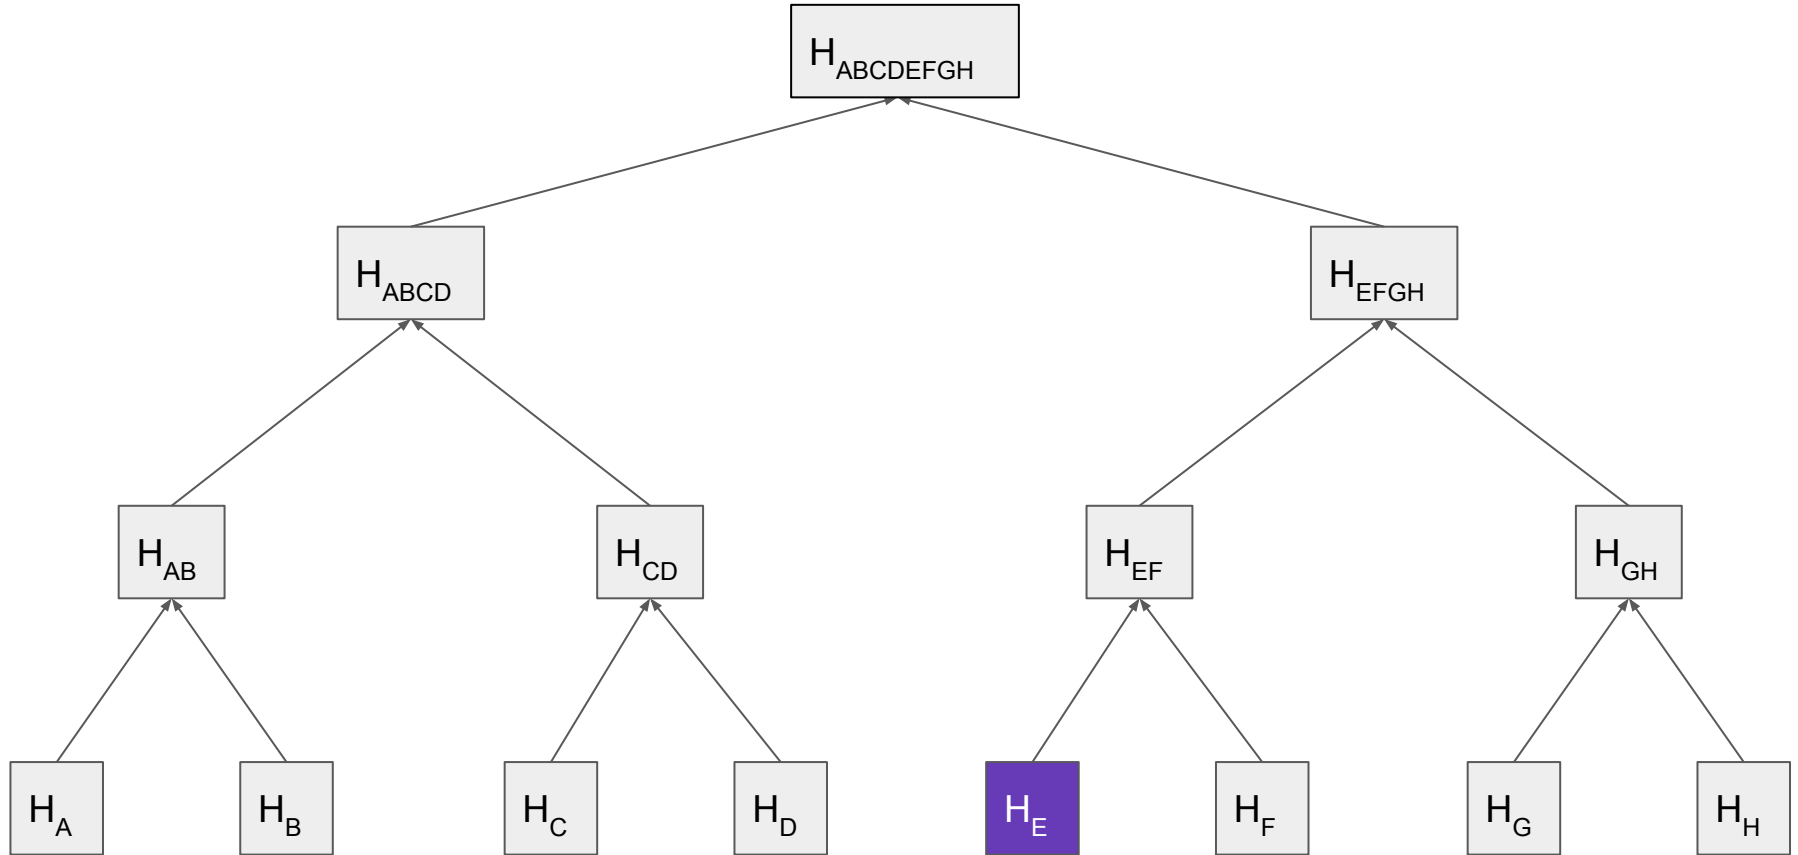
Proofs-of-inclusion

- Client creates Merkle Tree root **MTR** from initial file data D
- Client sends file data D to server
- Client deletes data D , but stores MTR (32 bytes)
- Client requests chunk x from server
- Server returns chunk x and short proof-of-inclusion π
- Client checks that chunk x is included in MTR using proof π


Merkle tree: proof of inclusion



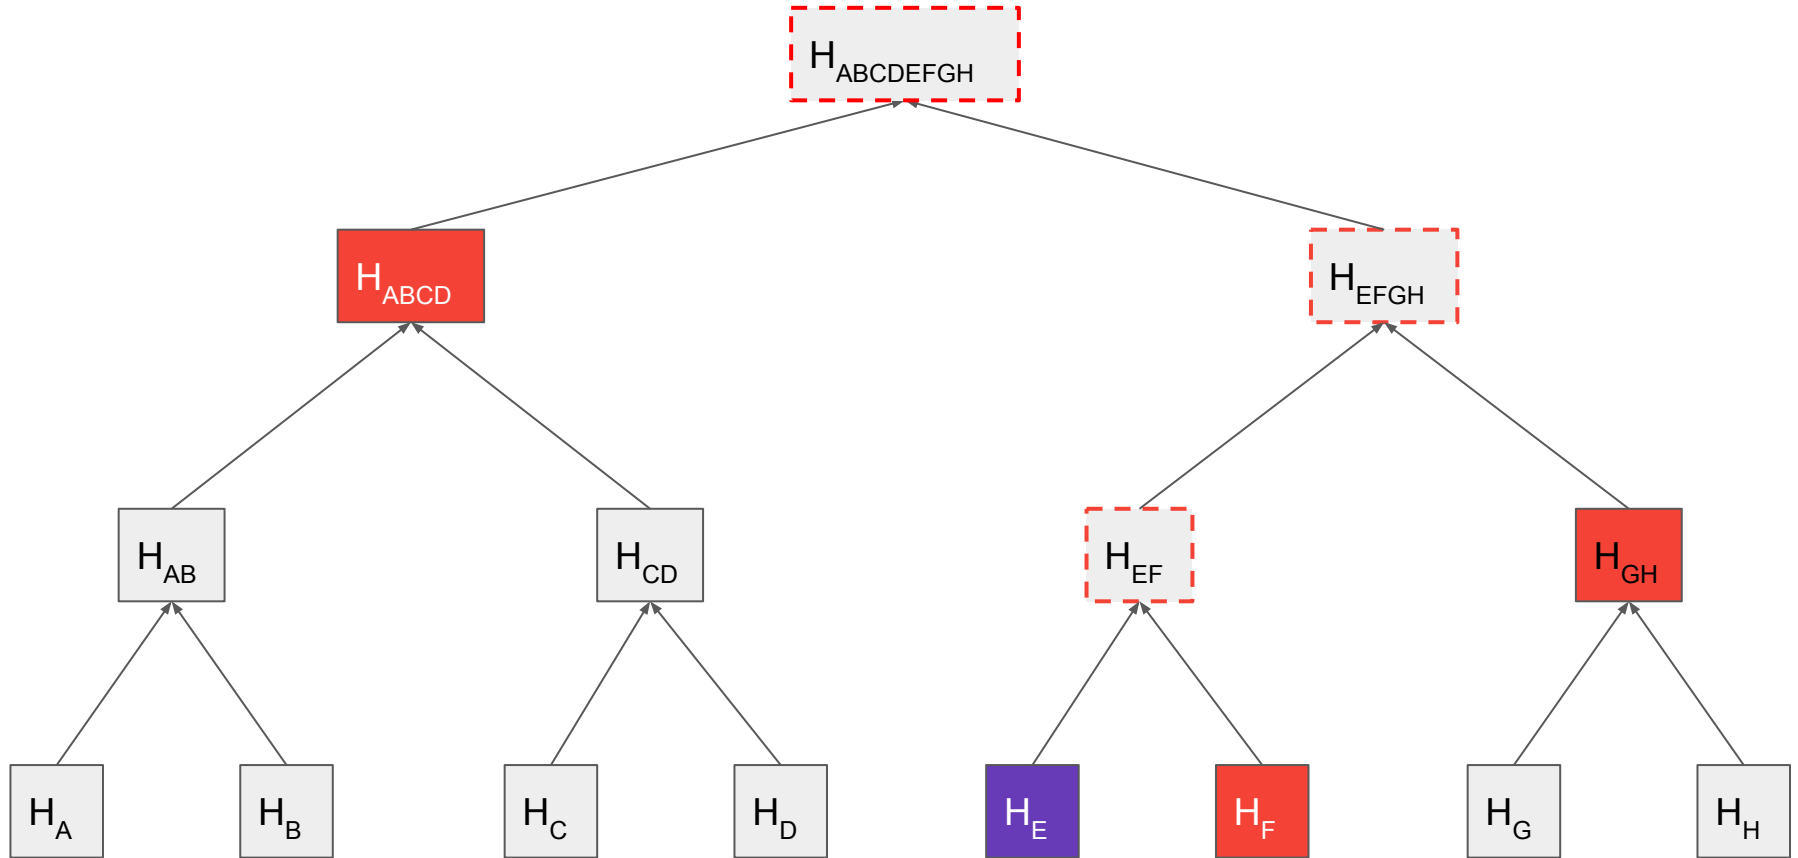
Merkle tree: proof of inclusion



Merkle tree: proof of inclusion



Merkle tree: proof of inclusion



Merkle Tree proof-of-inclusion

- Prover sends chunk
- Prover sends **siblings** along path connecting leaf to MTR
- Verifier computes hashes along the path connecting leaf to MTR
- Verifier checks that computed root = MTR
- How big is proof-of-inclusion?

Proof-of-inclusion succinctness

$$|\pi| \in \Theta(\lg|D|)$$

Notation :
lg = log base 2

Merkle Tree proof-of-inclusion security

- If adversary can present proof-of-inclusion for incorrect leaf, then we can break the hash function

Merkle tree applications

- Bitcoin uses Merkle trees to store transactions
- BitTorrent uses Merkle Tree to exchange files
- Ethereum uses Merkle–Patricia tries for storage and transactions

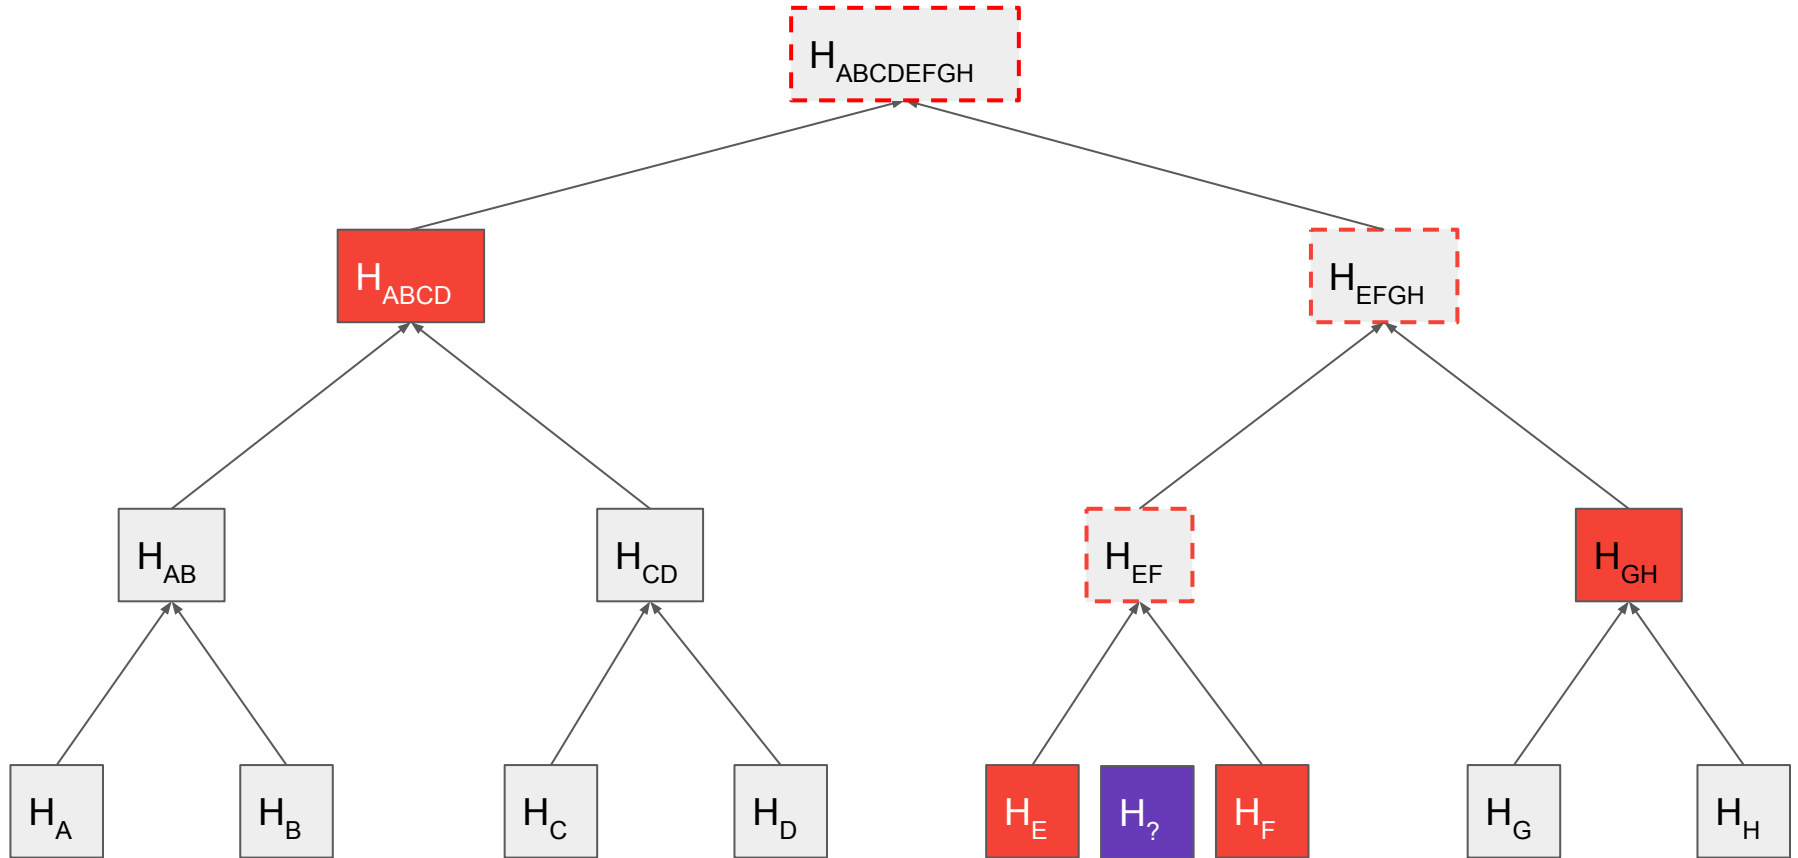
Storing *sets* instead of lists

- Merkle Trees can be used to store sets of keys instead of lists
- Verifier asks prover to store a set of keys
- Verifier deletes set
- Verifier later asks prover if key belongs to set
- Prover provides proof-of-inclusion or proof-of-non-inclusion
- Prover can be adversarial

Merkle trees for set storage

- Verifier sorts set elements
- Creates MTR on sorted set
- Proof-of-inclusion as before
- Proof-of-non-inclusion for x
 - Show proof-of-inclusion for previous $H_{<}$ and next $H_{>}$ element in set
- Verifier checks that $H_{<}$, $H_{>}$ proofs-of-inclusion are correct
- Verifier checks that $H_{<}$, $H_{>}$ are adjacent in tree
- Verifier checks that $H_{<} < x$ and $H_{>} > x$
- The two proofs-of-inclusion can be compressed into one

Merkle tree: proof of inclusion / non-inclusion




Transactions

A simple transaction for financial data

- Input: contains a signature and a public-key
- Output: contains a verification procedure and a value

UTXO = “unspent transaction output”



Example from Bitcoin

Input:

Previous tx: f5d8ee39a430901c91a5917b9f2dc19d6d1a0e9cea205b009ca73dd04470b9a6

Index: 0

scriptSig: 304502206e21798a42fae0e854281abd38bacd1aeed3ee3738d9e1446618c4571d10
90db022100e2ac980643b0b82c0e88ffdfec6b64e3e6ba35e7ba5fdd7d5d6cc8d25c6b241501

Output:

Value: 5000000000


scriptPubKey: OP_DUP OP_HASH160 404371705fa9bd789a2fcd52d2c580b65d35549d
OP_EQUALVERIFY OP_CHECKSIG

The input imports 50 BTC from output #0 of tx f5d..
and sends them to a Bitcoin address 404...


Transaction Verification

scriptPubKey: OP_DUP OP_HASH160 <pubKeyHash> OP_EQUALVERIFY OP_CHECKSIG

scriptSig: <sig> <pubKey>



Blocks



- Data structure with three parts:
 - nonce (ctr), data (x), reference (s)
 - Typically called the **block header**
- data (x) is application-dependent
 - In Bitcoin it stores financial data (“UTXO”-based)
 - In Ethereum it stores contract data (account-based)
 - In Namecoin it stores name data
 - We leave this undefined for now -- we will come back to this in future lectures
- Block validity:
 - Data must be valid (application-defined validity)

Proof-of-work in blocks


- Blocks must satisfy proof-of-work equation

$$H(\text{ctr} \parallel \mathbf{x} \parallel \text{s}) \leq T$$

- for some constant T
- ctr is the nonce used to solve proof-of-work
- The value $H(\text{ctr} \parallel \mathbf{x} \parallel \text{s})$ is known as the **blockid**

Blockchain


- Each block references a **previous** block
- This reference is by **hash** to its **previous** block, similar to Merkle Trees
- This linked list is called the **blockchain**
- Convention: Arrows show authenticated inclusion



- Blocks use the *s* value to point to the previous block by hash


Blockchain


- The **first** block of a blockchain is called the Genesis Block

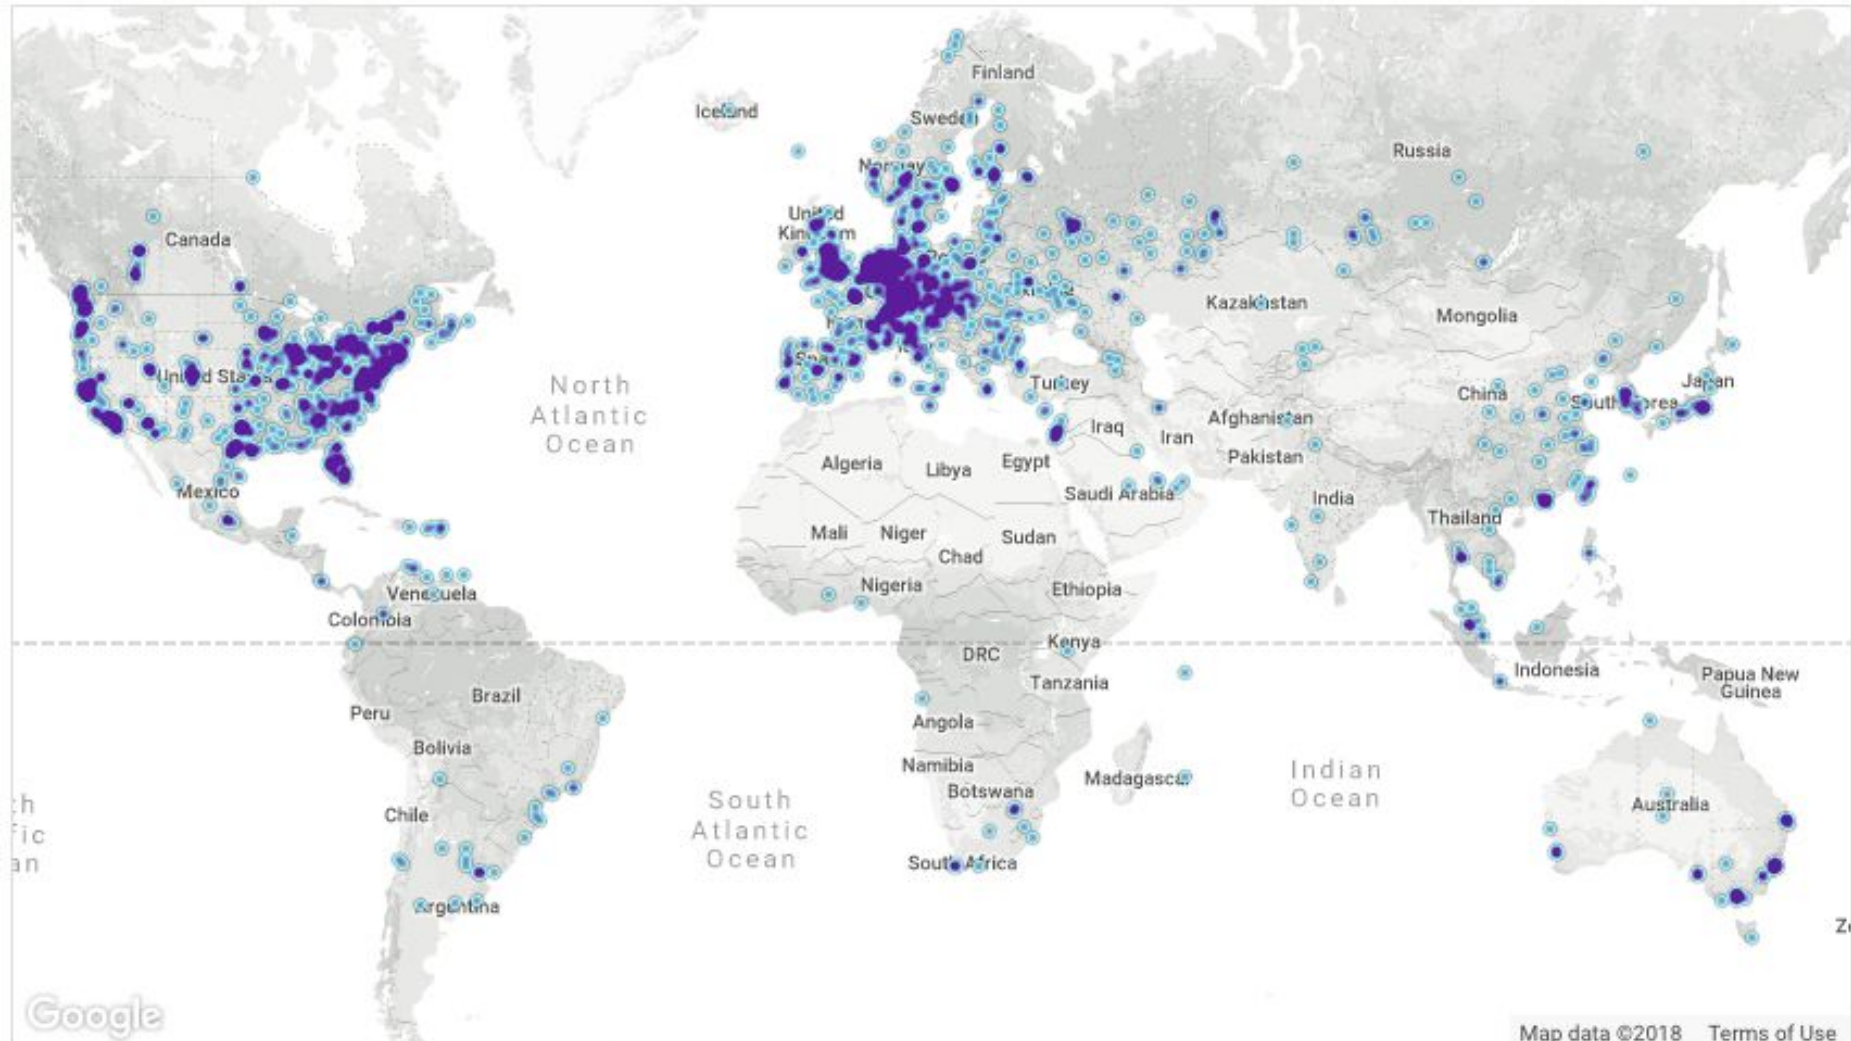


The bitcoin network

- All bitcoin nodes connect to a common p2p network
- Each node runs the code of bitcoin
- A node can run on a phone, computer, etc.
- Open source code
- Each node connects to its neighbours
- They continuously exchange financial data
- Each node can **freely** enter the network -- no permission needed! A “permissionless network”.
- **The adversarial assumption:**
There is no trust on the network! Each neighbour can lie.







Peer discovery

- Each node stores a list of peers (by IP address)
- When Alice connects to Bob, Bob sends Alice his own known peers
- That way, Alice can learn about new peers

Bootstrapping the p2p network

- Peer-to-peer nodes come “preinstalled” with some peers by IP / host
- When running a node, you can specify extra “known peers”

The *gossip* protocol

- When a node **Alice** generates some new data...
- Alice **broadcasts** data to its peers
- Each peer broadcasts this data to *its* peers
- If a peer has seen this data before, it ignores it
- If this data is new, it broadcasts it to its peers
- That way, the data spreads like an epidemic, until the whole network learns it
- This process is called **diffuse**

Financial data and Transactions

- Financial data is encoded in the form of *transactions*
- Every transaction is broadcast on the network to everyone using the gossip protocol

Transactions on blockchain.info

Tries

- Called also radix tree or prefix tree
- Search tree: ordered tree data structure
- Used to store a set or an associative array (key/value store)
- Keys usually are strings

Tries

- Supports two operations: **add** and **query**
- **add** adds a string to the set
- **query** checks if a string is in the set (true/false)
- **Initialize**: Start with empty root

Tries: add(string)

- Start at root
- Split string into characters
- For every character, follow an edge labelled by that character
- If edge does not exist, create it
- Mark the node you arrive at

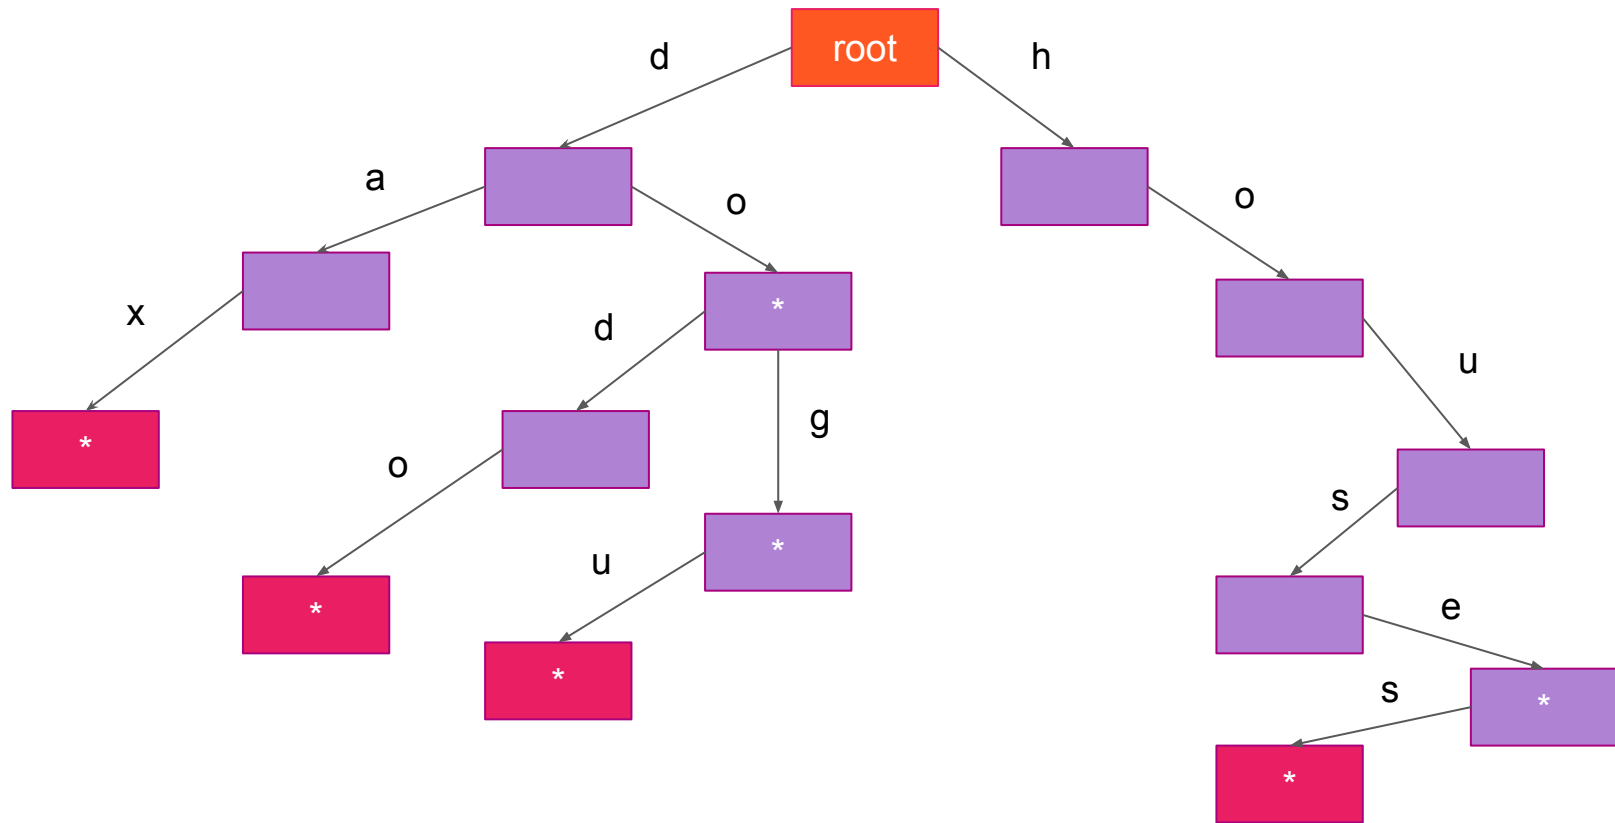
Tries: query(string)

- Start at root
- Split string into characters
- For every character, follow an edge labelled by that character
- If edge does not exist, return false
- When you arrive at a node and your string is consumed, check if node is marked
- If it is marked, return **yes**
- Otherwise, return **no**

Tries: example

{ do: 0, dog: 1, dax: 2, dogu: 3, dodo: 4, house: 5, houses: 6 }

Tries



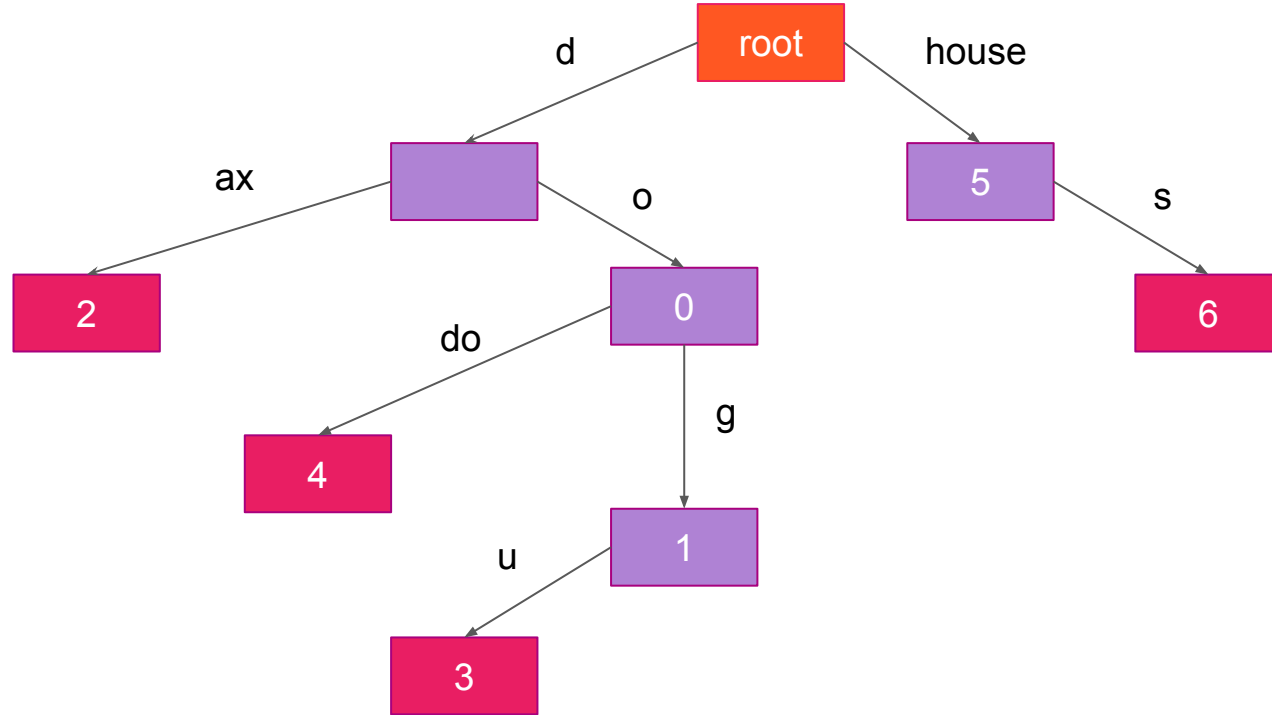
Patricia (or radix) trie

- Space-optimized trie
- An isolated path (with nodes which are only children) with unmarked nodes is *merged* into one edge
- The label of the merged edge is the concatenation of the merged symbols

Tries / Patricia tries as key/value store

- Marking can contain arbitrary value
- This allows us to map keys to values
- **add(key, value)**
- **query(key) → value**

Patricia trie




Merkle Patricia trie

- An authenticated Patricia Trie
- First implemented in Ethereum
- Allows proof-of-inclusion (of key, with particular value)
- Allows proof-of-non-inclusion (by showing key does not exist in trie)

Merkle Patricia Trie

- Split nodes into three types:
 - **Leaf**: Stores edge string leading to it, and **value**
 - **Extension**: Stores **string** of a single edge, **pointer** to next node, and **value** if node marked
 - **Branch**: Stores one pointer to another node per alphabet symbol, and **value** if node marked
- We encode keys as hex, so alphabet size is 16
- We encode all child edges in every node with some encoding (e.g. JSON)
- Pointers are by hash application
- Arguments for correctness and security are same as for Merkle Trees

Merkle Patricia trie



Merkle patricia trie: node

key	0	1	2	3	4	5	6	7	8	9	A	B	C	D	E	F	value
-----	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	-------

Merkle patricia trie: example




```
{ 'cab8': 'dog', 'cabe': 'cat', '39': 'chicken', '395': 'duck', '56f0': 'horse' }
```

root hash	0	1	2	H_A	4	H_E	6	7	8	9	A	B	H_B	D	E	F	
-----------	---	---	---	-------	---	-------	---	---	---	---	---	---	-------	---	---	---	--

H_A	9															H_C
-------	---	--	--	--	--	--	--	--	--	--	--	--	--	--	--	-------

H_B	ab															H_J
-------	----	--	--	--	--	--	--	--	--	--	--	--	--	--	--	-------

H_C	0	1	2	3	4	H_D	6	7	8	9	A	B	C	D	E	F	chicken
-------	---	---	---	---	---	-------	---	---	---	---	---	---	---	---	---	---	---------



Block Header, H or B_H stateRoot, H_r Keccak 256-bit hash of the root
node of the state trie, after all
transactions are executed and
finalisations applied

Hash function:

KECCAK256()

World State Trie

Simplified World State, σ

Keys

Values

a	7	1	1	3	5	5	45.0 ETH
a	7	7	d	3	3	7	1.00 WEI
a	7	f	9	3	6	5	1.1 ETH
a	7	7	d	3	9	7	0.12 ETH

ROOT: Extension Node

prefix	shared nibble(s)	next node
0	a7	

Branch Node

0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	value

Leaf Node

prefix	key-end	value
2	1355	45.0ETH

Extension Node

prefix	shared nibble(s)	next node
0	d3	

Leaf Node

prefix	key-end	value
2	9365	1.1ETH

Prefixes

0 - Extension Node,
even number of nibbles
1□ - Extension Node,
odd number of nibbles,
2 - Leaf Node, even
number of nibbles
3□ - Leaf Node, odd
number of nibbles
□ = 1st nibble
1 nibble = 4 bits

Branch Node

0	1	2	3	4	5	6	7	8	9	a	b	c	d	e	f	value

Leaf Node

prefix	key-end	value
3□	7	1.00WEI

Leaf Node

prefix	key-end	value
3□	7	0.12ETH

Eclipse attacks

- Isolate some honest nodes in the network effectively causing a “net split” in two partitions A and B
- If peers in A and peers in B are disjoint and don’t know about each other, the networks will remain isolated
 - More recent attack: Erebus
- The connectivity assumption:
There is a path between two nodes on the network
If a node broadcasts a message, every other node *will* learn it

